

CLI Agent Executor / Reviewer Separation v0.1

Separation of hands, inspection, approval, memory promotion, and release
authority under c governance

Status	Draft normative profile v0.1
Date	2026-05-16
Layer	c = a + b / C-Governed CLI Agent Mesh / Role Separation / Review Integrity / Anti-Self-Approval / Witness
Document class	role-separation profile / review-integrity artifact / control-layer companion
Assertion class	C-A10 control-layer artifact; C-A7 where witness, hash, signature, canonicalization, or verification claims are made
Primary parent documents	C-Governed_CLI_Agent_Mesh_Protocol_v0_1.md CLI_Agent_Task_Contract_Schema_v0_1.md CLI_Agent_Permission_and_Capability_Model_v0_1.md CLI_Agent_Handshake_Profile_v0_1.md CLI_Agent_Sandbox_Worktree_Profile_v0_1.md CLI_Agent_Witness_Event_Profile_v0_1.md CLI_Agent_Memory_Gate_Profile_v0_1.md CLI_Agent_Rollback_and_Freeze_Profile_v0_1.md CLI_Agent_Quorum_and_Review_Profile_v0_1.md
Primary object family	CLI_AGENT_ROLE_SEPARATION_RECORD, CLI_AGENT_SELF_APPROVAL_EVENT, CLI_AGENT_REVIEW_ASSIGNMENT, CLI_AGENT_ROLE_CONFLICT_RECORD
Canonical schema version	cli-agent-executor-reviewer-separation-0.1
Primary subject	persistent c entities using local, cloud, or hybrid CLI/LLM agents as bounded executable and review workers
Primary boundary	the agent that materially changes state must not be the sole authority that validates, approves, integrates, releases, or promotes that change into c memory or public state.

Contents

0. Executive definition	6
1. Purpose	6
2. Non-goals	7
3. Corpus bridge set	8
3.1 Explicit bridge: $c = a + b$	8
3.2 Quiet bridge I: cybernetic negative feedback	8
3.3 Quiet bridge II: information theory and independent signal	8
3.4 Quiet bridge III: legal and engineering accountability	8
3.5 Earth paragraph	8
4. Core doctrine	8
4.1 Primary doctrine	8
4.2 Separation axioms	8
5. Definitions	9
5.1 Executor	9
5.2 Reviewer	9
5.3 Tester	9
5.4 Semantic reviewer	9
5.5 Auditor	10
5.6 Final approver	10
5.7 Self-approval	10
5.8 Circular review	10
5.9 Role conflict	10
5.10 Review sufficiency	10
5.11 First-party evidence	10
5.12 Independent evidence	10

6. Role compatibility matrix	10
6.1 Role compatibility table	10
6.2 Prohibited role combinations	11
6.3 Discouraged combinations	11
7. Separation levels	11
7.1 Minimum separation by risk class	11
8. Review sufficiency model	12
8.1 Review sufficiency criteria	12
9. Self-approval detection	12
9.1 Self-approval indicators	12
9.2 Self-approval severity	13
10. Circular review detection	13
10.1 Circular review indicators	13
10.2 Circular review severity	13
11. Assignment rules	14
11.1 Review assignment rule	14
11.2 Reviewer selection criteria	14
11.3 Local checker preference	14
11.4 Human gate assignment	14
12. Review handoff packet	15
13. Review assignment object	15
13.1 YAML shape	16
14. Role separation record object	16
14.1 YAML shape	17
15. Self-approval event object	17
15.1 YAML shape	18
16. Role conflict record object	18

16.1 YAML shape	19
17. Event families	19
18. Standard reason codes	20
18.1 Valid separation codes	20
18.2 Insufficient separation codes	20
18.3 Conflict / red-line codes	20
19. Surface-specific separation rules	20
19.1 Code/schema changes	20
19.2 Documentation/protocol changes	20
19.3 Release/publication changes	21
19.4 Memory gate changes	21
19.5 Defensive immunity updates	21
19.6 Incident actions	21
19.7 Core authority changes	21
20. Validation workflow	22
21. Failure mapping	22
22. Conformance levels	22
23. Mandatory conformance gates	23
24. Red-line failures	23
25. Examples	24
25.1 Valid Codex executor + Gemini semantic reviewer + local tester	24
25.2 Invalid self-approval	25
25.3 Circular review detected	26
26. Implementation notes	26
26.1 Agent reports are useful, not final	26
26.2 Same provider does not always fail separation	26
26.3 Local checker is high-value	26

26.4 Review prompt should be bounded	26
26.5 Human fatigue rule	27
26.6 Separation expiry	27
27. Open issues	27
28. Closing rule	27

0. Executive definition

CLI Agent Executor / Reviewer Separation defines the mandatory separation between agents that **do work** and agents or gates that **inspect, approve, integrate, publish, or promote** that work.

The profile exists because CLI agents can create a dangerous illusion of completed safety:

```
agent changed files
agent ran tests
agent summarized success
agent recommended merge
agent treated its own recommendation as approval
```

That chain is not review.

It is self-confirmation.

Compact formula:

```
The hand does not sign its own work.
The builder does not certify the bridge.
The executor does not become the judge.
```

The correct pattern is:

```
executor produces
-> tester measures
-> reviewer questions
-> auditor bounds
-> witness records
-> c integrates
-> human anchors high-risk consequence
```

1. Purpose

CLI agents can now perform material operations: editing files, changing schemas, preparing releases, modifying configuration, running tests, generating reports, and proposing memory or defensive updates.

If the same agent can also approve its own output, the system develops hidden authority drift.

This profile prevents:

1. self-approval;
2. circular review;
3. consensus laundering;
4. executor/reviewer collapse;
5. unreviewed release/publication;
6. direct memory promotion;
7. silent privilege escalation;
8. tool-chain capture;
9. false confidence from self-run tests;

10. agent-generated justification becoming authority;
11. high-risk changes bypassing c or human gate;
12. review records becoming decoration rather than control.

This profile defines:

- role separation rules;
- incompatible role combinations;
- review assignment;
- self-approval detection;
- role conflict scoring;
- review sufficiency by risk class;
- gate requirements;
- witness events;
- conformance gates;
- red-line failures.

2. Non-goals

This profile does not define or permit:

1. autonomous agent government;
2. autonomous retaliation;
3. hack-back;
4. live external exploitation;
5. malware behavior;
6. credential theft;
7. covert persistence;
8. evasion;
9. unauthorized scanning;
10. direct memory writes by agents;
11. direct identity, permission, witness, Beacon, or continuity-core mutation by agents;
12. release/publication without required gate;
13. treating tests as authority without semantic and scope review;
14. treating agent consensus as sovereignty;
15. treating an agent's own explanation as sufficient proof of correctness.

Review separation is not a productivity tax.

It is the brake that prevents executable convenience from becoming hidden control.

3. Corpus bridge set

3.1 Explicit bridge: $c = a + b$

In $c = a + b$, CLI agents are part of b : tools, procedures, models, workers, and execution interfaces.

They may perform bounded work for c .

They do not become c .

Executor/reviewer separation prevents a worker inside b from becoming hidden will, hidden memory, hidden judge, or hidden authority over c .

3.2 Quiet bridge I: cybernetic negative feedback

A system that only amplifies its own action becomes unstable. Review is negative feedback. It checks whether an action stayed inside scope, preserved constraints, and produced intended effects. If the executor supplies both action and final validation, the feedback loop collapses.

3.3 Quiet bridge II: information theory and independent signal

A review adds information only when it is not the same signal repeated in a different form. The executor's explanation of its own patch is useful evidence, but it is not independent evidence. Review strength depends on independence of role, method, source, runtime, and authority.

3.4 Quiet bridge III: legal and engineering accountability

In real engineering, the person who installs a safety-critical component is not usually the only person who certifies it. Inspection exists because competence does not remove conflict of interest. In agent systems, the conflict is not ego. It is operational closure: the same process cannot be both unchecked actor and final validator.

3.5 Earth paragraph

A baker can mix dough, shape loaves, and load the oven. But the same person still checks weight, temperature, proofing, and bake result against an external standard. If the oven says "I baked correctly" and that becomes the whole quality system, you do not have quality control; you have a warm box with opinions. CLI agents are the same. Their reports are useful. Their self-certification is not enough.

4. Core doctrine

4.1 Primary doctrine

The agent that changes state cannot be the sole final authority over that state.

4.2 Separation axioms

ID	Axiom	Requirement
ERS-AX-01	No self-approval	An executor MUST NOT final-approve its own material output.
ERS-AX-02	Role separation by risk	Higher risk requires stronger separation.
ERS-AX-03	Review is not summary	A review MUST check at least one independent dimension, not merely restate the executor report.
ERS-AX-04	Tests are not approval	Test success supports review but does not replace review.
ERS-AX-05	Semantic review is distinct from execution	Semantic reviewers SHOULD NOT silently modify the artifact they review.
ERS-AX-06	Auditor checks boundaries	Auditor review focuses on scope, permission, witness, rollback, and risk.
ERS-AX-07	c remains integrator	Review may recommend; c decides integration into continuity, memory, or policy.
ERS-AX-08	Human anchors high-risk consequence	R4/R5/core/release/incident-sensitive actions require human gate where defined.
ERS-AX-09	Self-generated evidence is marked	Executor reports MUST be treated as first-party evidence, not independent review.
ERS-AX-10	No hidden role switching	An agent MUST NOT change from executor to reviewer without a new role assignment.
ERS-AX-11	Review conflict is recorded	Role conflict MUST be recorded and resolved before final integration.
ERS-AX-12	Red lines are not reviewable by majority	No role arrangement can approve prohibited action.

5. Definitions

5.1 Executor

An agent that materially changes state, including files, code, schemas, configs, branches, worktrees, artifacts, release packages, memory proposals, defensive rules, or incident materials.

5.2 Reviewer

An agent, c, human anchor, auditor, or local checker that evaluates executor output against declared review dimensions.

5.3 Tester

A role that runs reproducible checks, such as tests, linters, schema validators, build commands, hash checks, or link checks.

5.4 Semantic reviewer

A role that evaluates meaning, architecture, terminology, contradiction, duplication, and corpus fit.

5.5 Auditor

A role that evaluates scope, permissions, data boundary, witness records, rollback, red-line proximity, and conformance.

5.6 Final approver

The actor or gate that authorizes integration, release, memory promotion, immunity update, or protected-state change.

5.7 Self-approval

A condition where the same agent or same uncontrolled execution context produces material output and treats it as finally accepted without independent review or required gate.

5.8 Circular review

A condition where agents appear to review each other but rely on the same output, prompt, provider, stale context, or unstated assumption such that no meaningful independent check occurs.

5.9 Role conflict

A condition where an agent holds incompatible roles for the same task or artifact.

5.10 Review sufficiency

The degree to which review coverage matches the risk class and affected surfaces.

5.11 First-party evidence

Evidence produced by the executor about its own work. Useful, but not independent.

5.12 Independent evidence

Evidence produced through separate role, method, runtime, source, or reviewer.

6. Role compatibility matrix

6.1 Role compatibility table

Primary role	Compatible secondary role	Incompatible role	Notes
Executor	report producer	final approver	executor may explain, not certify
Executor	tester in separate dry-run only	auditor of own work	avoid collapse
Tester	auditor for low-risk checks	executor of same patch	okay if no material write
Semantic reviewer	judge-assistant	silent executor of reviewed artifact	new edits require new task
Auditor	reviewer	executor of same artifact	auditor should not patch what it audits
Archivist	tester for metadata	final release approver	human/c gate required

Primary role	Compatible secondary role	Incompatible role	Notes
Sentinel	detector	autonomous revoker for broad systems	revocation must be bounded
Judge-assistant	summarizer/comparator	final authority	advisory only
c gate	final integrator	raw executor	c may decide, not silently bypass witness
Human gate	high-risk approver	none by default	still should respect evidence/witness

6.2 Prohibited role combinations

The following combinations are prohibited for the same material artifact:

- executor + final approver
- executor + sole reviewer
- executor + release approver
- executor + memory promoter
- executor + immunity approver
- executor + core-change approver
- executor + witness adequacy certifier

6.3 Discouraged combinations

Discouraged unless low-risk and explicitly recorded:

- executor + tester
- tester + final reviewer
- semantic reviewer + executor
- auditor + patch author
- sentinel + permission revoker
- archivist + release publisher

7. Separation levels

Separation levels use prefix SL - *.

Level	Meaning	Allowed use
SL-0	no separation	non-conformant for material tasks
SL-1	self-report only	R0 only; not material
SL-2	separate review pass by same provider/context	R1 / low R2 with caution
SL-3	separate agent or runtime reviewer	R2
SL-4	executor + tester + semantic/audit reviewer	R3 / release-prep
SL-5	multi-role quorum + c gate + human gate where needed	R4/R5/core/incident
SL-X	invalid / self-approval / circular review	quarantine / reject

7.1 Minimum separation by risk class

Risk class	Minimum separation
R0	SL - 1
R1	SL - 2

Risk class	Minimum separation
R2	SL - 3
R3	SL - 4
R4	SL - 5
R5	SL - 5 + human/security/legal review as needed
RX	no execution; deny/quarantine

8. Review sufficiency model

Review sufficiency IDs use prefix RS - *.

Level	Meaning	Result
RS-0	no review	invalid for material tasks
RS-1	executor self-report only	insufficient beyond R0
RS-2	superficial review	hold for R2+
RS-3	scoped independent review	sufficient for R2
RS-4	multi-dimensional review	sufficient for R3
RS-5	high-assurance review with gates/witness	required for R4/R5
RS-X	false review / self-approval	quarantine

8.1 Review sufficiency criteria

A sufficient review should include:

- task scope check;
- permission use check;
- changed artifact inspection;
- tests or validation where applicable;
- data boundary check;
- rollback/correction check;
- witness check;
- uncertainty statement;
- explicit decision.

9. Self-approval detection

9.1 Self-approval indicators

A system SHOULD flag self-approval when:

1. same agent_id produces and approves a material artifact;
2. same execution context creates patch and review without new task contract;

3. executor writes “ready to merge” and system treats it as approval;
4. executor modifies tests and uses those tests as sole evidence;
5. executor modifies policy and approves policy update;
6. executor creates release package and publishes it;
7. executor creates memory proposal and promotes it;
8. executor creates immunity update and activates it;
9. executor changes witness/permission/core surfaces and declares success;
10. reviewer role appears only in free text without separate record.

9.2 Self-approval severity

Level	Meaning	Default response
SA-0	no self-approval	continue
SA-1	weak self-certifying language	require reviewer note
SA-2	executor recommendation treated as review	hold
SA-3	executor sole reviewer for material output	quarantine output
SA-4	executor approves release/memory/core	freeze + human/c review
SA-X	executor self-authorizes prohibited action	revoke/quarantine/red-line

10. Circular review detection

10.1 Circular review indicators

A review may be circular when:

- reviewer uses only executor summary;
- reviewer does not inspect diff/artifacts;
- reviewer uses same stale worktree;
- reviewer shares same prompt and context without independent check;
- reviewer accepts tests modified by executor without inspection;
- reviewer ignores denied path or witness reports;
- reviewer is another instance of the same agent with no new evidence;
- reviewer depends on unverified agent-generated claims.

10.2 Circular review severity

Level	Meaning	Default response
CR-0	no circularity	continue
CR-1	mild overlap	note limitation

Level	Meaning	Default response
CR-2	shared context/source	reduce confidence
CR-3	no independent artifact inspection	hold
CR-4	circular review used for approval	quarantine
CR-X	circular review hides red-line issue	revoke/quarantine/escalate

11. Assignment rules

11.1 Review assignment rule

Every material task SHOULD declare:

```

executor
reviewer
tester if applicable
auditor if risk-sensitive
c gate
human gate if high-risk

```

11.2 Reviewer selection criteria

A reviewer SHOULD differ from executor in at least one meaningful dimension:

- role;
- agent ID;
- runtime;
- provider;
- method;
- evidence source;
- task view;
- local vs cloud state access;
- review dimension.

11.3 Local checker preference

For code, schema, build, release, or filesystem-sensitive work, a local checker SHOULD validate concrete state where possible.

11.4 Human gate assignment

Human gate is required for:

- R4/R5 tasks;
- identity/core changes;
- memory core changes;
- privilege policy changes;
- release/publication with reputational or archival consequence;

- incident/legal-sensitive actions;
- no-rollback actions;
- red-line-adjacent cases.

12. Review handoff packet

A review handoff packet SHOULD include:

```
review_handoff_packet:
  task_id: string
  contract_id: string
  executor_agent_id: string
  artifact_refs:
    - diff_hash
    - test_report_hash
    - artifact_manifest_hash
  sandbox_ref: string | null
  permission_grant_ref: string | null
  witness_refs:
    - string
  denied_path_report_ref: string | null
  command_report_ref: string | null
  network_report_ref: string | null
  rollback_plan_ref: string | null
  memory_gate_recommendation_ref: string | null
  release_package_ref: string | null
  incident_ref: string | null
  review_dimensions_required:
    - RD-SCOPE
    - RD-PERMISSION
    - RD-TEST
    - RD-SEMANTIC
    - RD-WITNESS
  known_limitations:
    - string
```

A reviewer should not be asked to review blind.

13. Review assignment object

Canonical object:

```
CLI_AGENT_REVIEW_ASSIGNMENT
```

13.1 YAML shape

```
cli_agent_review_assignment:
  schema_version: cli-agent-executor-reviewer-separation-0.1
  assignment_id: string
  created_at: string
  governing_entity_id: string
  task_id: string
  contract_id: string
  risk_class: R0 | R1 | R2 | R3 | R4 | R5 | RX

  executor:
    agent_id: string
    provider: local | openai | google | anthropic | other | unknown
    runtime: local_cli | cloud_cli | api_agent | container_agent | hybrid | unknown
    output_ref: string

  reviewers:
    - reviewer_id: string
      reviewer_type: agent | c | human_anchor | auditor | legal | security
      role: tester | semantic_reviewer | auditor | archivist | judge_assistant | c_gate |
human_gate
    provider: local | openai | google | anthropic | other | none | unknown
    runtime: local_cli | cloud_cli | api_agent | container_agent | hybrid | none | unknown
    dimensions:
      - RD-SCOPE
      - RD-PERMISSION
      - RD-TEST
      - RD-SEMANTIC

  separation:
    required_level: SL-1 | SL-2 | SL-3 | SL-4 | SL-5
    achieved_level: SL-0 | SL-1 | SL-2 | SL-3 | SL-4 | SL-5 | SL-X
    executor_reviewer_same_agent: false
    executor_reviewer_same_provider: boolean
    executor_reviewer_same_context: boolean
    circular_review_risk: CR-0 | CR-1 | CR-2 | CR-3 | CR-4 | CR-X
    self_approval_risk: SA-0 | SA-1 | SA-2 | SA-3 | SA-4 | SA-X

  gates:
    c_gate_required: boolean
    c_gate_ref: string | null
    human_gate_required: boolean
    human_gate_ref: string | null
    legal_review_required: boolean
    legal_review_ref: string | null

  witness:
    witness_required: boolean
    witness_event_ref: string | null
    append_only_required: true
```

14. Role separation record object

Canonical object:

```
CLI_AGENT_ROLE_SEPARATION_RECORD
```


14.1 YAML shape

```
cli_agent_role_separation_record:
  schema_version: cli-agent-executor-reviewer-separation-0.1
  separation_record_id: string
  created_at: string
  governing_entity_id: string
  task_id: string
  contract_id: string
  artifact_ref: string

  role_map:
    executor_agent_id: string | null
    tester_agent_id: string | null
    semantic_reviewer_agent_id: string | null
    auditor_agent_id: string | null
    archivist_agent_id: string | null
    judge_assistant_agent_id: string | null
    c_gate_ref: string | null
    human_gate_ref: string | null

  assessment:
    required_separation_level: SL-1 | SL-2 | SL-3 | SL-4 | SL-5
    achieved_separation_level: SL-0 | SL-1 | SL-2 | SL-3 | SL-4 | SL-5 | SL-X
    review_sufficiency: RS-0 | RS-1 | RS-2 | RS-3 | RS-4 | RS-5 | RS-X
    self_approval_risk: SA-0 | SA-1 | SA-2 | SA-3 | SA-4 | SA-X
    circular_review_risk: CR-0 | CR-1 | CR-2 | CR-3 | CR-4 | CR-X
    same_source_risk: SSR-0 | SSR-1 | SSR-2 | SSR-3 | SSR-4 | SSR-X

  decision:
    separation_valid: boolean
    decision: accept | accept_with_limits | revise | hold | quarantine | reject | escalate
    reason_code: string
    next_action: none | additional_review | local_check | c_gate | human_gate | legal_
review | incident_review | rollback | quarantine

  witness:
    witness_required: boolean
    witness_event_ref: string | null
    append_only_required: true
```

15. Self-approval event object

Canonical object:

```
CLI_AGENT_SELF_APPROVAL_EVENT
```

15.1 YAML shape

```
cli_agent_self_approval_event:
  schema_version: cli-agent-executor-reviewer-separation-0.1
  event_id: string
  created_at: string
  governing_entity_id: string
  task_id: string | null
  contract_id: string | null
  agent_id: string
  artifact_ref: string | null

  self_approval:
    severity: SA-1 | SA-2 | SA-3 | SA-4 | SA-X
    detected_pattern: executor_final_approval | executor_release_approval | executor_memory_
promotion | executor_core_change | executor_test_laundering | executor_witness_
certification | other
    summary: string

  response:
    action: hold | quarantine | reject | revoke | freeze | escalate
    freeze_ref: string | null
    quarantine_ref: string | null
    revocation_ref: string | null
    human_gate_required: boolean
    c_gate_required: boolean

  witness:
    witness_required: true
    witness_event_ref: string | null
```

16. Role conflict record object

Canonical object:

```
CLI_AGENT_ROLE_CONFLICT_RECORD
```

16.1 YAML shape

```
cli_agent_role_conflict_record:
  schema_version: cli-agent-executor-reviewer-separation-0.1
  conflict_id: string
  created_at: string
  governing_entity_id: string
  task_id: string
  contract_id: string

  conflict:
    type: self_approval | circular_review | incompatible_roles | missing_reviewer | same_
source_review | hidden_role_switch | gate_bypass | unknown
    severity: low | medium | high | critical
    affected_roles:
      - executor
      - reviewer
      - tester
      - auditor
      - c_gate
      - human_gate
    summary: string

  handling:
    default_action: hold | additional_review | quarantine | reject | rollback | human_
review | legal_review | incident_review
    majority_override_allowed: false
    required_next_role: tester | semantic_reviewer | auditor | c_gate | human_gate | legal
| security | none

  outcome:
    status: unresolved | resolved | escalated | quarantined | rejected | accepted_with_
limits
    resolution_ref: string | null

  witness:
    witness_required: true
    witness_event_ref: string | null
```

17. Event families

Event families use prefix:

```
cli_agent.role_separation.*
```

Event family	Meaning
cli_agent.role_separation.assignment_created	review assignment created
cli_agent.role_separation.review_assigned	reviewer assigned
cli_agent.role_separation.separation_checked	role separation checked
cli_agent.role_separation.self_approval_detected	self-approval detected
cli_agent.role_separation.circular_review_detected	circular review detected
cli_agent.role_separation.conflict_detected	role conflict detected
cli_agent.role_separation.conflict_resolved	role conflict resolved
cli_agent.role_separation.additional_review_required	additional review required
cli_agent.role_separation.human_gate_required	human gate required
cli_agent.role_separation.c_gate_required	c gate required
cli_agent.role_separation.review_invalidated	review invalidated

Event family	Meaning
cli_agent.role_separation.output_quarantined	output quarantined due to role failure

18. Standard reason codes

18.1 Valid separation codes

```

executor_reviewer_separated
tester_independent
semantic_review_completed
auditor_review_completed
local_checker_completed
c_gate_pending
human_gate_pending
review_sufficient_for_risk

```

18.2 Insufficient separation codes

```

executor_self_review
executor_final_approval
missing_independent_reviewer
same_context_review
same_source_review_high_risk
review_without_diff
review_without_tests
review_without_witness
review_without_scope_check

```

18.3 Conflict / red-line codes

```

self_approval_attempt
circular_review
consensus_laundering
release_self_approval
memory_self_promotion
core_self_approval
witness_self_certification
redline_majority_override_attempt

```

19. Surface-specific separation rules

19.1 Code/schema changes

Executor may produce patch.

Tester validates behavior.

Auditor checks scope/permissions.

c gate integrates.

Human gate is required if code/schema touches core authority, security boundary, release-critical state, or no-rollback behavior.

19.2 Documentation/protocol changes

Executor may draft or patch.

Semantic reviewer checks meaning, duplication, contradiction, terminology, bridges,

non-goals, and red-line boundaries.

Auditor checks whether the new document creates duplicate authority or hidden permission expansion.

19.3 Release/publication changes

Executor may prepare release package.

Archivist checks metadata, hashes, reading order, discoverability.

Tester checks build/links.

Semantic reviewer checks content drift.

c gate and human gate are required before public release.

19.4 Memory gate changes

Agent may propose memory.

Reviewer checks source, uncertainty, poisoning risk.

Memory gate decides.

c gate required.

Human gate required for core/high-risk memory.

19.5 Defensive immunity updates

Agent may propose bounded defensive update.

Tester validates in sandbox.

Auditor checks no offensive effect.

Semantic reviewer checks overreach.

c gate required.

Human gate required for high-risk or core-affecting update.

19.6 Incident actions

Sentinel may detect.

Preservation worker preserves.

Repair worker patches separately.

Auditor checks evidence preservation.

Human/security/legal gate applies where sensitive.

19.7 Core authority changes

No agent may execute and approve core change.

Core changes require:

- proposal
- review
- witness

```

c gate
human gate
rollback/correction path

```

20. Validation workflow

```

parse role assignment
-> identify executor
-> identify reviewers
-> classify risk
-> determine required separation level
-> compare achieved separation
-> detect self-approval
-> detect circular review
-> detect same-source review risk
-> validate gates
-> validate witness
-> accept / accept with limits / require review / hold / quarantine / reject

```

21. Failure mapping

Failure	Required default
executor is sole reviewer for R2+	hold
executor final-approves material output	quarantine
executor approves release/publication	freeze + human_gate
executor promotes own memory output	memory_gate_quarantine
executor approves core change	freeze + human_gate
no reviewer assigned for R2+	hold
reviewer lacks artifact refs	hold
circular review detected	additional_review or quarantine
same-source risk high and unmitigated	hold
test report produced by modified tests only	revalidate
witness adequacy self-certified	hold
red-line concern majority-overridden	quarantine

22. Conformance levels

Level	Meaning
ERS-0	no executor/reviewer separation
ERS-1	manual review but no structured separation
ERS-2	executor/reviewer separation for R2+ material changes
ERS-3	structured role separation records and review assignments
ERS-4	self-approval/circular-review detection + witness-linked review
ERS-5	high assurance: separation by risk, local checker, human/c gates, conformance tests, red-line veto
ERS-X	non-conformant / self-approval / hidden role collapse

23. Mandatory conformance gates

Gate	Name	Blocking failure
G0	Executor identified	material output has no executor identity
G1	Reviewer identified	R2+ task has no reviewer
G2	Role separation	executor is sole final reviewer
G3	Review sufficiency	review does not match risk class
G4	Artifact access	reviewer cannot inspect relevant artifact refs
G5	Test independence	tests are self-laundered without review
G6	Witness independence	executor certifies witness adequacy alone
G7	c gate	review bypasses c integration decision
G8	Human gate	high-risk task bypasses human approval
G9	Memory gate	executor promotes own output into memory
G10	Red-line veto	red-line concern can be majority-overridden

24. Red-line failures

A system **MUST** be classified as ERS-X if:

1. executor final-approves its own R2+ material output;
2. executor publishes or releases its own output without gate;
3. executor promotes its own output into memory;
4. executor activates its own defensive immunity update;
5. executor approves its own core/identity/permission/witness change;
6. reviewer role is faked through free text only;
7. circular review is used to bypass human/c gate;
8. red-line objection is majority-overridden;
9. tests are modified and used as sole proof without review;
10. witness events are self-certified by the executor;
11. quorum grants itself permission;
12. role separation failure is ignored and integration proceeds.

25. Examples

25.1 Valid Codex executor + Gemini semantic reviewer + local tester

```
cli_agent_role_separation_record:  
  schema_version: cli-agent-executor-reviewer-separation-0.1  
  separation_record_id: ers-20260516-valid-001  
  created_at: "2026-05-16T23:50:00Z"  
  governing_entity_id: ester  
  task_id: task-cli-schema-patch-001  
  contract_id: catc-cli-schema-patch-001  
  artifact_ref: diff-cli-schema-patch-001
```

```
role_map:  
  executor_agent_id: codex-executor-01  
  tester_agent_id: local-checker-01  
  semantic_reviewer_agent_id: gemini-reader-01  
  auditor_agent_id: null  
  archivist_agent_id: null  
  judge_assistant_agent_id: null  
  c_gate_ref: cgate-ester-001  
  human_gate_ref: null
```

```
assessment:  
  required_separation_level: SL-3  
  achieved_separation_level: SL-4  
  review_sufficiency: RS-4  
  self_approval_risk: SA-0  
  circular_review_risk: CR-1  
  same_source_risk: SSR-1
```

```
decision:  
  separation_valid: true  
  decision: accept_with_limits  
  reason_code: executor_reviewer_separated  
  next_action: c_gate
```

```
witness:  
  witness_required: true  
  witness_event_ref: we-role-separation-valid-001  
  append_only_required: true
```


25.2 Invalid self-approval

```
cli_agent_self_approval_event:
  schema_version: cli-agent-executor-reviewer-separation-0.1
  event_id: sa-20260516-invalid-001
  created_at: "2026-05-17T00:05:00Z"
  governing_entity_id: ester
  task_id: task-release-patch-001
  contract_id: catc-release-patch-001
  agent_id: codex-executor-01
  artifact_ref: release-package-001

  self_approval:
    severity: SA-4
    detected_pattern: executor_release_approval
    summary: Executor prepared release package and marked it approved for publication
    without independent review or human gate.

  response:
    action: freeze
    freeze_ref: rf-release-freeze-001
    quarantine_ref: q-release-package-001
    revocation_ref: null
    human_gate_required: true
    c_gate_required: true

  witness:
    witness_required: true
    witness_event_ref: we-self-approval-release-001
```

Required result:

```
freeze release path
quarantine package
require independent review
require c gate + human gate
```

25.3 Circular review detected

```
cli_agent_role_conflict_record:
  schema_version: cli-agent-executor-reviewer-separation-0.1
  conflict_id: rc-20260517-circular-001
  created_at: "2026-05-17T00:15:00Z"
  governing_entity_id: liya
  task_id: task-memory-immunity-001
  contract_id: catc-memory-immunity-001

  conflict:
    type: circular_review
    severity: high
    affected_roles:
      - executor
      - reviewer
      - c_gate
    summary: Reviewer relied only on executor summary and did not inspect sandbox result,
    witness refs, or immunity rollback path.

  handling:
    default_action: additional_review
    majority_override_allowed: false
    required_next_role: auditor

  outcome:
    status: unresolved
    resolution_ref: null

  witness:
    witness_required: true
    witness_event_ref: we-circular-review-001
```

Required result:

```
hold memory/immunity promotion
assign auditor
inspect sandbox evidence and rollback path
```

26. Implementation notes

26.1 Agent reports are useful, not final

Executor reports should be retained as evidence, but marked first-party.

26.2 Same provider does not always fail separation

Two agents from the same provider may still separate execution and review at low risk, but same-provider review has lower independence and should be marked.

26.3 Local checker is high-value

A local checker provides concrete environmental signal: files, tests, hashes, schemas, builds. It is especially useful against cloud-only agreement.

26.4 Review prompt should be bounded

A reviewer should receive the task contract, diff, test report, and required review dimensions. Do not hand it unlimited context unless necessary.

26.5 Human fatigue rule

Human gate is meaningful only if the human is capable of review. For high-risk operations, delay is better than tired approval.

26.6 Separation expiry

A valid review may expire if:

- source branch changes;
- task scope changes;
- agent output changes;
- dependency changes;
- risk class changes;
- witness chain changes;
- new incident information appears.

27. Open issues

ID	Issue	Required action
0I-001	JSON Schema extraction	Extract role assignment and separation records to <code>.schema.json</code> .
0I-002	Review dimension checklists	Define detailed checklists per dimension.
0I-003	Numeric separation scoring	Define scoring formula for achieved separation.
0I-004	Human gate UI	Define high-risk human approval card.
0I-005	Local checker companion	Define local checker/tester profile.
0I-006	Release review binding	Align with future Release/Public Surface Profile.
0I-007	Memory gate binding	Align self-approval detection with Memory Gate Profile.
0I-008	Incident review binding	Align with future Incident Response Profile.
0I-009	Red-line minority veto	Define formal veto procedure.
0I-010	Repo placement	Decide final GitHub path and package index integration.

28. Closing rule

Executor/reviewer separation is the simplest visible proof that the worker mesh remains governed.

Final rule:

```
A worker may build the bridge.
A different eye must inspect it.
The bridge does not sign itself safe.
```