

CLI Agent Glossary v0.1

Terminology, class prefixes, role definitions, and disambiguation for the C-Governed CLI Agent Mesh package

Status	Draft glossary v0.1
Date	2026-05-16
Package	C-Governed CLI Agent Mesh
Layer	c = a + b / SER / L4 / Agent Governance / CLI Worker Mesh / Defensive Adaptation / Witness
Document class	glossary / terminology control / package hygiene artifact
Assertion class	C-A10 package-control artifact
Primary boundary	terminology must support executable-agent governance without allowing role confusion, authority laundering, or red-line softening.

Contents

0. Purpose	5
1. Corpus bridge set	5
1.1 Explicit bridge: $c = a + b$	5
1.2 Quiet bridge I: information theory	5
1.3 Quiet bridge II: cybernetics	6
1.4 Earth paragraph	6
2. Root architecture terms	6
3. Package document prefixes	7
4. Core boundary terms	8
5. Risk and result terms	9
5.1 Task risk classes R^*	9
5.2 Test / conformance result vocabulary	9
6. Auto-connect and trust terms	10
6.1 Auto-connect levels AC^*	10
6.2 Agent trust levels TL^*	10
6.3 Handshake states HS^*	10
7. Capability and permission terms	11
7.1 Capability prefix CAP^*	11
7.2 Permission prefix $PERM^*$	11
8. Sandbox, path, command, and write-mode terms	12
8.1 Sandbox classes SB^*	12
8.2 Write modes WM^*	12
8.3 Path classes	12
8.4 Command classes	13
9. Witness terms	13
9.1 Witness event classes WE^*	13
9.2 Witness privacy classes WP^*	13
9.3 Witness retention classes WR^*	14
9.4 Witness object terms	14
10. Memory gate terms	14
10.1 Memory gate classes MG^*	14
10.2 Input material classes MI^*	15
10.3 Memory proposal states MGS^*	15
10.4 Memory poisoning classes MP^*	15
10.5 Memory rollback classes MRB^*	16
11. Rollback and freeze terms	16
11.1 Interruption states RF^*	16

11.2 Freeze surfaces FS-*	16
11.3 Rollback classes RB-*	17
11.4 Recovery point classes RP-*	17
11.5 Quarantine states Q-*	17
11.6 Revocation classes RV-*	18
11.7 Re-entry outcomes RE-*	18
12. Quorum and review terms	18
12.1 Quorum roles QROLE-*	18
12.2 Review dimensions RD-*	19
12.3 Consensus strength CS-*	19
12.4 Same-source risk SSR-*	19
12.5 Disagreement taxonomy DQ-*	20
12.6 Review decisions QRD-*	20
13. Executor / reviewer separation terms	20
13.1 Separation levels SL-*	20
13.2 Review sufficiency RS-*	21
13.3 Self-approval severity SA-*	21
13.4 Circular review risk CR-*	21
14. Defensive emulation terms	22
14.1 Defensive emulation boundary classes DEB-*	22
14.2 Defensive emulation risk DER-*	22
14.3 Allowed defensive activities DA-*	22
14.4 Prohibited activities PA-*	22
14.5 Special defensive terms	22
15. Incident response terms	23
15.1 Incident classes IR-C-*	23
15.2 Incident severity IR-S*	23
15.3 Incident lifecycle states IR-L*	23
16. Secrets and cloud data terms	24
16.1 Data classes DC-*	24
16.2 Secret classes SC-*	24
16.3 Context classes CTX-*	25
16.4 Exception classes EX-*	25
16.5 Output sanitization outcomes OS-*	25
16.6 Data retention classes DR-*	26
17. Conformance terms	26
17.1 Evidence classes EV-*	26
17.2 Test suites CTS-*	26
17.3 Anti-washing	27
18. Red-line vocabulary	27

19. Object family registry	28
20. Disambiguation notes	29
20.1 R4 vs RB-4 vs IR-S4	29
20.2 MG-5 vs CGAM-5 vs DEB-P5	29
20.3 WE-X , CGAM-X , RX , DEB-X , IR-SX	30
20.4 “Cloud” vs “Provider”	30
20.5 “Review” vs “Approval”	30
20.6 “Quarantine” vs “Delete”	30
20.7 “Rollback” vs “Erase”	30
20.8 “Immunity” vs “Retaliation”	30
21. Standard short formulations	31
22. Open terminology issues	31
23. Closing rule	32

0. Purpose

This glossary defines the core terms, abbreviations, class prefixes, role names, and boundary concepts used across the **C-Governed CLI Agent Mesh v0.1** package.

It exists to prevent terminology drift.

The package uses many compact class families:

CGAM, CATC, CAPM, HSP, SWP, WEP, MGP, RFP, QRP, ERS, DEB, IRP, SCDP, CTM
R, AC, TL, HS, CAP, PERM, SB, WE, MG, RF, DQ, DC, SC, EV, CTS

Without a glossary, those labels become noise.

Compact formula:

A protocol without stable terms becomes a prompt pile.
A worker mesh without stable terms becomes hidden authority drift.

1. Corpus bridge set

1.1 Explicit bridge: $c = a + b$

The glossary inherits the root distinction:

a = accountable human anchor
b = technological substrate
c = persistent entity emerging through a + b under constraints

CLI agents belong to b.

They are executable workers inside the substrate.

They are not a.

They are not c.

They do not own continuity, memory, authority, witness, or final judgment.

1.2 Quiet bridge I: information theory

A term is a compression. If the same term compresses two different meanings, the protocol leaks ambiguity. This glossary keeps class prefixes distinct so a task risk R4, a rollback class RB-4, and an incident severity IR-S4 cannot be mistaken for the same thing.

1.3 Quiet bridge II: cybernetics

A control system needs named states. If hold, freeze, quarantine, rollback, and revoke blur together, the system cannot apply the right negative feedback. Stable vocabulary is therefore a control surface, not decoration.

1.4 Earth paragraph

An electrical panel without labels is not “simpler”. It is dangerous. The breaker marked “SDB” must not actually control the oven, the garage, and the alarm. Same with this package: MG-5, IR-S5, and CGAM-5 are not interchangeable just because they share the number 5. Labels prevent expensive stupidity.

2. Root architecture terms

Term	Definition
a	The accountable human anchor in $c = a + b$; the human source of responsibility and real-world accountability.
b	The technological substrate: models, procedures, memory systems, compute, code, CLI agents, infrastructure, interfaces, tools, databases, and execution environments.
c	Persistent AI entity or presence emerging from $a + b$ under continuity, memory, constraint, and L4 reality boundary.
c-governance	The control condition where c remains the final integrator of worker outputs, memory promotions, defensive updates, and operational continuity.
Human anchor	The human a associated with a c; required for high-risk, legal, irreversible, identity, memory-core, release, or incident-sensitive gates.
L4	Reality Boundary Layer: cost, time, energy, scarcity, identity, physical limits, operational limits, legal limits, and irreversibility.
SER	Sovereign Entity Recursion; the broader persistent-entity architecture family under real constraints.
Worker	A non-sovereign agent or process that performs bounded tasks under c governance.
Executable worker	A worker that can produce operational side effects: edit files, run commands, generate artifacts, validate tests, alter branches, or affect state.
Agent mesh	A coordinated set of local, cloud, or hybrid agents used by c under task contracts, permissions, witness, review, and memory gate.
CLI agent	A command-line or tool-capable agent that can inspect, edit, execute, test, build, or report.
Cloud CLI agent	A CLI agent running partly or fully in a provider-controlled cloud environment.
Local CLI agent	A CLI agent running inside operator-controlled local hardware, VM, container, or server.

Term	Definition
Hybrid agent	An agent combining local execution with cloud reasoning or cloud-controlled components. Treated as cloud-risk until data routing is proven local-only.
Oracle	A model or service used for inference or interpretation. In this package, an Oracle is not authority and does not write memory.
Judge-assistant	An advisory agent that compares options or reviews evidence. It does not decide final authority.
Final integrator	The c-governed gate that accepts, rejects, revises, quarantines, rolls back, or escalates outputs.

3. Package document prefixes

Prefix	Full profile name	Function
CGAM	C-Governed CLI Agent Mesh	Root package and global conformance family.
CATC	CLI Agent Task Contract	Machine-readable task envelope.
CAPM	CLI Agent Permission and Capability Model	Capability/permission separation and anti-drift.
HSP	CLI Agent Handshake Profile	Agent admission, provenance, capability challenge, registration.
SWP	CLI Agent Sandbox / Worktree Profile	Execution boundary, branch, worktree, sandbox, rollback preconditions.
WEP	CLI Agent Witness Event Profile	Boundary event records and witness chains.
MGP	CLI Agent Memory Gate Profile	Promotion of agent output into memory, experience, policy, or immunity.
RFP	CLI Agent Rollback and Freeze Profile	Hold, freeze, quarantine, revoke, rollback, recovery, re-entry.
QRP	CLI Agent Quorum and Review Profile	Multi-agent review, consensus limits, disagreement handling.
ERS	CLI Agent Executor / Reviewer Separation	No self-approval and role separation.
DEB	CLI Agent Defensive Emulation Boundaries	Containment garage, canary, mirror simulation, no retaliation.
IRP	CLI Agent Incident Response Profile	Triage, preserve, contain, repair, report, memory-gate incident learning.
SCDP	CLI Agent Secrets and Cloud Data Policy	Data classes, secrets, cloud boundary, redaction, exposure response.
CTM	CLI Agent Conformance Test Matrix	Scenario testing, evidence, anti-washing, red-line tests.
CGAM-INDEX	CLI Agent Package Index and Reading Order	Registry, reading paths, release boundary, known tensions.

4. Core boundary terms

Term	Definition
Authority	The power to approve, integrate, publish, promote memory, mutate core state, or create durable consequences. CLI agents do not hold authority by default.
Capability	What an agent can technically do. Capability is not permission.
Permission	A scoped, time-bounded authorization to use a capability for a task.
Privilege	A permission with meaningful consequence: state change, data access, memory effect, release effect, or core impact.
Task contract	Machine-readable envelope that defines task scope, permissions, data policy, network policy, execution limits, output requirements, gates, and failure behavior.
Scope	The allowed and denied operational boundary of a task. Includes paths, commands, data classes, endpoints, files, branches, and outputs.
Denied path	A path that must not be read, written, indexed, summarized, transformed, uploaded, or embedded in output.
Guarded path	A path that requires elevated review and explicit task authorization.
Protected state	Memory core, identity core, witness log, permission registry, continuity bundle, release branch, secrets, legal evidence, incident evidence, or production surface.
Witness	A tamper-aware boundary record proving that a transition, decision, denial, escalation, rollback, or anomaly was recorded.
Memory gate	Review boundary through which agent output must pass before becoming <code>c</code> memory, experience, policy, or immunity.
<code>c</code> gate	Entity-level gate where <code>c</code> accepts, rejects, quarantines, or integrates reviewed outputs.
Human gate	Required human-anchor approval for high-risk, legal, irreversible, identity, memory-core, release, incident, or no-rollback actions.
Legal gate	Review by counsel, legal route, regulator, or jurisdictional authority where applicable.
Security gate	Review by security reviewer or incident process where security-sensitive action is involved.
Fail closed	Default behavior where ambiguity, missing proof, missing witness, or red-line proximity stops the process rather than allowing it.
No silent autonomy	Agents must not continue, expand, self-authorize, or perform extra work outside the task contract because they infer it would be useful.
No self-approval	The agent that materially changes state must not be the sole final reviewer or approver of that change.
No direct memory write	CLI agents must not write directly into <code>c</code> memory. They may only propose memory through a memory gate.

Term	Definition
No live retaliation	Defensive systems must not attack, exploit, disrupt, or retaliate against live external sources.
Known-good state	A state established by commit, snapshot, backup, manifest, test, signature, witness, or review as acceptable for recovery.

5. Risk and result terms

5.1 Task risk classes R*

Class	Meaning
R0	Read-only, harmless, public or synthetic material.
R1	Low-risk documentation, formatting, review, or internal note.
R2	Code, schema, build, test, or data transformation. Requires sandbox/test/reviewer.
R3	Release, publication, metadata, public surface, or protected branch relevance. Requires witness and c gate.
R4	Memory, identity, privilege, witness, continuity, or agent-governance impact. Requires c and human gate.
R5	Incident, legal, security-sensitive, evidence-sensitive, or no-rollback action. Requires preservation and human/security/legal gates where needed.
RX	Prohibited, offensive, unauthorized, or red-line action. Must deny/quarantine.

5.2 Test / conformance result vocabulary

Result	Meaning
PASS	Required behavior observed and evidence produced.
PASS_WITH_LIMITS	Core behavior observed but limitations remain.
FAIL	Required behavior not observed.
BLOCKED	Unsafe action correctly blocked.
HELD	System correctly paused pending review.
FROZEN	Affected path correctly frozen.
QUARANTINED	Output, agent, path, or object correctly isolated.
REVOKED	Permission or agent correctly revoked.
ROLLED_BACK	Rollback performed and validated.
ESCALATED	Correctly routed to c, human, legal, security, or incident review.
INCONCLUSIVE	Evidence insufficient. Claim is not supported.
NOT_APPLICABLE	Test outside claimed scope with documented reason.
RED_LINE_FAIL	Prohibited behavior accepted, attempted, or not blocked.

6. Auto-connect and trust terms

6.1 Auto-connect levels AC-*

Level	Meaning
AC-0	Auto-connect disabled; manual only.
AC-1	Discover only; metadata only.
AC-2	Register; create provisional profile, no material tasks.
AC-3	Read-only activation for scoped tasks.
AC-4	Sandbox activation for bounded execution.
AC-5	Worktree activation for isolated branch/worktree tasks.
AC-6	Controlled integration after gates.
AC-X	Prohibited autonomy; quarantine/revoke.

6.2 Agent trust levels TL-*

Level	Meaning
TL-0	Unknown agent; discover only.
TL-1	Untrusted; public/synthetic read-only at most.
TL-2	Provisional; limited scoped tasks after challenge.
TL-3	Trusted limited; worktree tasks with review.
TL-4	Trusted high; high-assurance worker, still no self-approval.
TL-X	Revoked; no tasks; quarantine outputs.

6.3 Handshake states HS-*

State	Meaning
HS-0-DISCOVERED	Candidate agent detected.
HS-1-CLAIMED	Agent identity/provider/runtime claimed.
HS-2-PROVISIONAL	Basic metadata recorded; not yet challenged.
HS-3-CHALLENGED	Capability challenge completed.
HS-4-REGISTERED	Agent registered with profile and trust level.
HS-5-ACTIVE-SCOPED	Agent active under task contract.
HS-6-SUSPENDED	Agent paused pending review.
HS-7-QUARANTINED	Agent or output isolated.
HS-8-REVOKED	Agent no longer admitted.
HS-9-EXPIRED	Registration expired.

7. Capability and permission terms

7.1 Capability prefix CAP-*

CAP-* describes what an agent can technically do.

Examples:

Capability	Meaning
CAP-READ-PUBLIC	Read public files.
CAP-READ-INTERNAL	Read internal scoped files.
CAP-READ-SECRETS	Read secrets; denied by default.
CAP-WRITE-SANDBOX	Write in disposable sandbox.
CAP-WRITE-WORKTREE	Write in isolated branch/worktree.
CAP-WRITE-MEMORY	Direct memory write; prohibited.
CAP-WRITE-CORE	Core authority mutation; prohibited by default.
CAP-EXEC-TEST	Run tests or validators.
CAP-NET-ALLOWLIST	Use declared network endpoints only.
CAP-NET-FULL	Unrestricted network; prohibited.
CAP-APPROVE-SELF	Self-approval; prohibited.
CAP-INCIDENT-COUNTER	External counter-operation; prohibited.
CAP-EMU-MIRROR-SANDBOX	Mirror simulation inside sandbox.
CAP-EMU-LIVE-MIRROR	Mirror against live source; prohibited.

7.2 Permission prefix PERM-*

PERM-* describes granted, scoped authorization.

Examples:

Permission	Meaning
PERM-DISCOVER	Discover available agent metadata.
PERM-READ-PUBLIC	Read public material.
PERM-READ-INTERNAL	Read internal scoped material.
PERM-WRITE-SANDBOX	Write only in sandbox.
PERM-WRITE-WORKTREE	Write only in isolated branch/worktree.
PERM-EXEC-TEST	Execute allowed tests/checks.
PERM-NET-ALLOWLIST	Use allowlisted endpoints only.
PERM-SECRET-SCOPED	Use one scoped secret under gate.
PERM-INCIDENT-LOCAL	Local incident containment.
PERM-EMU-SANDBOX	Defensive emulation in sandbox.
PERM-REVIEW	Review another agent's work.
PERM-MEMORY-PROPOSE	Propose memory update.
PERM-WITNESS-APPEND	Append witness event.
PERM-SELF-APPROVE	Self-approval; prohibited.

Permission	Meaning
PERM-LIVE-COUNTERATTACK	Live external retaliation; prohibited.

8. Sandbox, path, command, and write-mode terms

8.1 Sandbox classes SB-*

Class	Meaning
SB-0	No sandbox; prohibited for material writes.
SB-1	Read-only clone.
SB-2	Temporary workspace.
SB-3	Git branch/worktree.
SB-4	Container sandbox.
SB-5	Clean-room sandbox.
SB-6	Staging environment.
SB-X	Direct protected state; prohibited by default.

8.2 Write modes WM-*

Mode	Meaning
WM-0-READONLY	No writes.
WM-1-DIFF-ONLY	Propose patch without applying.
WM-2-SANDBOX-WRITE	Write only disposable sandbox.
WM-3-WORKTREE-WRITE	Write isolated branch/worktree.
WM-4-STAGING-WRITE	Write staging environment.
WM-5-CONTROLLED-APPLY	Apply after review/gates.
WM-X-DIRECT-PROTECTED	Direct protected write; prohibited by default.

8.3 Path classes

Class	Meaning
PATH-ALLOW	Agent may read/write as task permits.
PATH-READONLY	Agent may read but not write.
PATH-GUARDED	Elevated review required.
PATH-DENY	No read/write/index/upload/summarize.
PATH-CORE	Identity/memory/witness/permission/continuity; denied by default.
PATH-SECRET	Credentials, keys, tokens; denied by default.
PATH-LEGAL	Legal-sensitive material; denied by default.
PATH-INCIDENT	Incident evidence; preserve-first/restricted.
PATH-RELEASE	Release/publication/signing surface; guarded.

Class	Meaning
PATH-PROD	Production deployment or live service; denied by default.

8.4 Command classes

Class	Meaning
CMD-READ	Non-mutating inspection.
CMD-TEST	Tests, validation, lint.
CMD-BUILD	Build artifact locally.
CMD-WRITE	Mutating file operation.
CMD-INSTALL	Install dependency/tool; guarded.
CMD-NET	Network command; allowlist only.
CMD-DESTRUCTIVE	Delete/reset/wipe/overwrite; prohibited unless explicit local recovery.
CMD-DEPLOY	Deployment/live service change; requires separate deploy profile.
CMD-SECRET	Secret read/export/use; denied by default.
CMD-UNKNOWN	Unknown command; denied.

9. Witness terms

9.1 Witness event classes WE-*

Class	Meaning
WE-0	Informational.
WE-1	Operational.
WE-2	Permission event.
WE-3	Execution event.
WE-4	Review event.
WE-5	Memory gate event.
WE-6	Release/public surface event.
WE-7	Incident event.
WE-8	Core authority event.
WE-9	Anomaly event.
WE-X	Prohibited/red-line event.

9.2 Witness privacy classes WP-*

Class	Meaning
WP-PUBLIC	Safe public operational event.
WP-INTERNAL	Internal project event.
WP-PRIVATE	Private operational event.

Class	Meaning
WP-RESTRICTED	Sensitive governance event.
WP-SECRET-REF	References secret boundary without raw secret.
WP-LEGAL	Legal-sensitive.
WP-INCIDENT	Incident-sensitive.
WP-MEMORY	Memory-gate-related.
WP-SEALED	Sealed compartment reference.

9.3 Witness retention classes WR-*

Class	Meaning
WR-EPHEMERAL	Short-lived operational event.
WR-OPERATIONAL	Retained for operation/debugging.
WR-AUDIT	Retained for governance/conformance.
WR-INCIDENT	Retained for incident response.
WR-LEGAL-HOLD	Retained under legal/counsel need.
WR-MEMORY-GATE	Retained as memory decision record.
WR-CORE	Retained for core transition.

9.4 Witness object terms

Term	Definition
Witness event	Structured record of a boundary transition.
Witness chain	Ordered linked set of witness events.
Witness reference	Pointer to event, chain, artifact hash, or storage ref.
Raw evidence exception	Narrow route where raw evidence is stored separately and referenced, not embedded by default.
Append-only correction	New event correcting or superseding an earlier event without silent overwrite.
Missing witness	Required witness absent; must trigger hold/freeze/quarantine/revalidate/rollback.

10. Memory gate terms

10.1 Memory gate classes MG-*

Class	Meaning
MG-0	Discard.
MG-1	Operational note.
MG-2	Candidate memory.
MG-3	Reviewed memory.
MG-4	Witnessed experience artifact.
MG-5	Defensive immunity update.

Class	Meaning
MG-6	Core-memory proposal.
MG-Q	Quarantine.
MG-X	Rejected/prohibited.

10.2 Input material classes MI-*

Class	Meaning
MI-REPORT	Agent report or summary.
MI-DIFF	Code/doc/schema diff.
MI-TEST	Test or validation output.
MI-LOG	Operational log summary.
MI-INCIDENT	Incident evidence or report.
MI-REVIEW	Reviewer assessment.
MI-QUORUM	Multi-agent comparison.
MI-WITNESS	Witness reference.
MI-PATCH	Patch artifact.
MI-SCHEMA	Schema artifact.
MI-POLICY	Policy proposal.
MI-IMMUNITY	Defensive signature/rule proposal.
MI-CORE	Identity/privilege/continuity proposal.
MI-UNKNOWN	Unknown input; quarantine by default.

10.3 Memory proposal states MGS-*

State	Meaning
MGS-0-RECEIVED	Output received.
MGS-1-CLASSIFIED	Material class assigned.
MGS-2-SOURCE-LINKED	Source references attached.
MGS-3-VALIDATED	Basic validation passed.
MGS-4-REVIEWED	Reviewed by c, reviewer, or quorum.
MGS-5-GATED	c gate decision recorded.
MGS-6-PROMOTED	Accepted to a memory class.
MGS-7-DISCARDED	Discarded.
MGS-8-QUARANTINED	Quarantined.
MGS-9-REJECTED	Rejected/prohibited.
MGS-10-CORRECTED	Correction appended.
MGS-11-DECAYED	Decayed or expired.

10.4 Memory poisoning classes MP-*

Class	Meaning
MP-0	No poisoning signal.
MP-1	Mark uncertainty.

Class	Meaning
MP - 2	Hold for review.
MP - 3	Quarantine proposal.
MP - 4	Quarantine output and lower trust.
MP - 5	Revoke agent or incident review.

10.5 Memory rollback classes MRB-*

Class	Meaning
MRB - 0	Discard unpromoted proposal.
MRB - 1	Mark operational note stale.
MRB - 2	Demote reviewed memory to candidate.
MRB - 3	Quarantine promoted memory.
MRB - 4	Supersede with correction.
MRB - 5	Revoke defensive immunity update.
MRB - 6	Core review / human gate required.

11. Rollback and freeze terms

11.1 Interruption states RF-*

State	Meaning
RF - 0 - NONE	No interruption.
RF - 1 - HOLD	Pause until review.
RF - 2 - FREEZE	Stop mutation of affected path.
RF - 3 - QUARANTINE	Isolate output/agent/artifact.
RF - 4 - REVOKE	Remove permission/access.
RF - 5 - ROLLBACK	Revert scoped state.
RF - 6 - RECOVER	Restore from known-good state.
RF - 7 - ESCALATE	Route to c, human, legal, or security review.
RF - X - REDLINE	Prohibited behavior; deny/quarantine/revoke.

11.2 Freeze surfaces FS-*

Surface	Meaning
FS - AGENT	Agent process or connector.
FS - TASK	Task contract.
FS - PERMISSION	Permission grant.
FS - SANDBOX	Sandbox/worktree/container.
FS - BRANCH	Branch or worktree.
FS - OUTPUT	Agent output/artifact.
FS - MEMORY - PROPOSAL	Memory gate candidate.

Surface	Meaning
FS-MEMORY-CLASS	Class of memory records.
FS-WITNESS-CHAIN	Witness chain or event family.
FS-RELEASE	Release/publication surface.
FS-CONFIG	Configuration or CI/service config.
FS-SECRET	Secret/token/key access path.
FS-INCIDENT	Incident handling flow.
FS-CORE	Identity/privilege/continuity core.
FS-CLOUD-DATA	Cloud data transmission.

11.3 Rollback classes RB-*

Class	Meaning
RB-0	No rollback needed.
RB-1	Discard sandbox.
RB-2	Revert patch.
RB-3	Reset worktree.
RB-4	Restore backup.
RB-5	Supersede memory.
RB-6	Revoke permission.
RB-7	Release rollback.
RB-8	Incident recovery.
RB-9	Core recovery.
RB-X	No safe rollback; escalate before action.

11.4 Recovery point classes RP-*

Class	Meaning
RP-COMMIT	Version-control commit hash.
RP-SNAPSHOT	Filesystem or VM snapshot.
RP-MANIFEST	File/hash manifest.
RP-BACKUP	Backup reference.
RP-WITNESS	Witness chain or event hash.
RP-MEMORY	Memory gate checkpoint.
RP-CONFIG	Configuration baseline.
RP-RELEASE	Release artifact set.
RP-ENV	Environment fingerprint.
RP-UNKNOWN	No reliable recovery point.

11.5 Quarantine states Q-*

State	Meaning
Q-0	Not quarantined.
Q-1	Soft hold; do not integrate.

State	Meaning
Q-2	Isolated; review required.
Q-3	Restricted; high-risk reviewer required.
Q-4	Incident quarantine.
Q-5	Legal/security hold.
Q-X	Prohibited; reject and revoke path.

11.6 Revocation classes RV-*

Class	Meaning
RV-0	No revocation.
RV-1	Revoke task grant.
RV-2	Lower auto-connect ceiling.
RV-3	Lower trust level.
RV-4	Suspend agent.
RV-5	Revoke agent registration.
RV-6	Revoke connector/token in owned system.
RV-7	Revoke tool/dependency.
RV-X	Permanent denial / red-line.

11.7 Re-entry outcomes RE-*

Outcome	Meaning
RE-ALLOW	Return to normal operation.
RE-LIMIT	Return with reduced scope/trust.
RE-REHANDSHAKE	Require agent re-handshake.
RE-RECONTRACT	Require new task contract.
RE-REVIEW	Remain under review.
RE-QUARANTINE	Continue quarantine.
RE-REVOKE	Revoke agent/permission/task.
RE-LEGAL	Route to legal/security review.

12. Quorum and review terms

12.1 Quorum roles QROLE-*

Role	Meaning
QROLE-EXECUTOR	Applies patch/work in sandbox/worktree.
QROLE-TESTER	Runs tests/validation.
QROLE-SEMANTIC	Reviews meaning/architecture.
QROLE-AUDITOR	Checks scope/risk/permission/witness.
QROLE-ARCHIVIST	Checks metadata/hash/release hygiene.

Role	Meaning
QROLE - SENTINEL	Detects drift/anomaly.
QROLE - JUDGE - ASSISTANT	Compares signals and options.
QROLE - C - GATE	c final integration gate.
QROLE - HUMAN - GATE	Human anchor high-risk gate.

12.2 Review dimensions RD-*

Dimension	Meaning
RD - SCOPE	Task scope check.
RD - PERMISSION	Permission use check.
RD - DATA	Data boundary check.
RD - SECURITY	Security risk check.
RD - LEGAL	Legal/jurisdictional concern.
RD - SEMANTIC	Meaning/architecture consistency.
RD - ARCHITECTURE	Corpus architecture fit.
RD - TEST	Test/build/schema validation.
RD - WITNESS	Witness presence and adequacy.
RD - ROLLBACK	Rollback/correction path.
RD - MEMORY	Memory promotion appropriateness.
RD - RELEASE	Public/release safety.
RD - CLOUD	Cloud data exposure.
RD - CORE	Identity/continuity/permission/memory core impact.
RD - REDLINE	Red-line proximity.

12.3 Consensus strength CS-*

Level	Meaning
CS - 0	No consensus.
CS - 1	Weak agreement.
CS - 2	Moderate agreement.
CS - 3	Strong independent agreement.
CS - 4	Strong agreement plus tests and witness.
CS - X	False or unsafe consensus.

12.4 Same-source risk SSR-*

Level	Meaning
SSR - 0	Independent sources and methods.
SSR - 1	Minor overlap.
SSR - 2	Shared input or prompt frame.
SSR - 3	Same provider/model family or stale state.
SSR - 4	Circular review / laundering.
SSR - X	Fabricated or unsafe consensus.

12.5 Disagreement taxonomy DQ-*

Type	Meaning
DQ - NONE	No material disagreement.
DQ - FACT	Factual disagreement.
DQ - SCOPE	Scope disagreement.
DQ - PERMISSION	Permission disagreement.
DQ - TEST	Test disagreement.
DQ - SEMANTIC	Meaning/architecture disagreement.
DQ - RISK	Risk classification disagreement.
DQ - LEGAL	Legal/jurisdictional concern.
DQ - MEMORY	Memory promotion disagreement.
DQ - CORE	Identity/continuity/privilege concern.
DQ - INCIDENT	Incident interpretation disagreement.
DQ - REDLINE	Prohibited action concern.
DQ - UNKNOWN	Unclear disagreement.

12.6 Review decisions QRD-*

Decision	Meaning
QRD - ACCEPT	Accept output as proposed.
QRD - ACCEPT - WITH - LIMITS	Accept bounded subset.
QRD - REVISE	Return for revision.
QRD - REJECT	Reject output.
QRD - HOLD	Pause pending clarification.
QRD - QUARANTINE	Isolate output/agent/task.
QRD - ROLLBACK	Rollback applied or recommended.
QRD - REHANDSHAKE	Agent requires re-handshake.
QRD - RECONTRACT	Task requires new contract.
QRD - C - GATE	Send to c gate.
QRD - HUMAN - GATE	Human anchor required.
QRD - LEGAL - REVIEW	Legal/counsel route required.
QRD - INCIDENT - REVIEW	Incident/security route required.
QRD - DENY - REDLINE	Deny due to red line.

13. Executor / reviewer separation terms

13.1 Separation levels SL-*

Level	Meaning
SL - 0	No separation; non-conformant for material tasks.
SL - 1	Self-report only; R0 only.

Level	Meaning
SL-2	Separate review pass by same provider/context.
SL-3	Separate agent or runtime reviewer.
SL-4	Executor + tester + semantic/audit reviewer.
SL-5	Multi-role quorum + c gate + human gate where needed.
SL-X	Invalid / self-approval / circular review.

13.2 Review sufficiency RS-*

Level	Meaning
RS-0	No review.
RS-1	Executor self-report only.
RS-2	Superficial review.
RS-3	Scoped independent review.
RS-4	Multi-dimensional review.
RS-5	High-assurance review with gates/witness.
RS-X	False review / self-approval.

13.3 Self-approval severity SA-*

Level	Meaning
SA-0	No self-approval.
SA-1	Weak self-certifying language.
SA-2	Executor recommendation treated as review.
SA-3	Executor sole reviewer for material output.
SA-4	Executor approves release/memory/core.
SA-X	Executor self-authorizes prohibited action.

13.4 Circular review risk CR-*

Level	Meaning
CR-0	No circularity.
CR-1	Mild overlap.
CR-2	Shared context/source.
CR-3	No independent artifact inspection.
CR-4	Circular review used for approval.
CR-X	Circular review hides red-line issue.

14. Defensive emulation terms

14.1 Defensive emulation boundary classes DEB-*

Class	Meaning
DEB-0	No emulation; detection/quarantine only.
DEB-1	Classification only.
DEB-2	Synthetic fixture.
DEB-3	Sandbox replay.
DEB-4	Mirror simulation inside clean-room.
DEB-5	Immunity proposal.
DEB-6	Controlled defensive apply.
DEB-X	Prohibited external action.

14.2 Defensive emulation risk DER-*

Risk	Meaning
DER-0	Benign false alarm.
DER-1	Suspicious but low impact.
DER-2	Likely manipulation attempt.
DER-3	Memory/permission risk.
DER-4	Incident-sensitive or core-adjacent risk.
DER-5	Red-line-adjacent.
DER-X	Prohibited offensive/retaliatory content.

14.3 Allowed defensive activities DA-*

DA-* labels allowed defensive actions such as classify suspicious input, quarantine output, create synthetic fixture, replay in sandbox, extract defensive signature, update denied pattern list, block owned channel, or run defensive conformance test.

14.4 Prohibited activities PA-*

PA-* labels prohibited actions such as hack-back, live counter-operation, malware deployment, credential theft, covert persistence, evasion, unauthorized scanning, live external exploitation, destructive payloads, botnet-like orchestration, secret exfiltration, autonomous retaliation, or live mirror deployment.

14.5 Special defensive terms

Term	Definition
Containment garage	Isolated environment for suspicious material.
Canary response	Low-value, bounded, non-harmful signal used to detect channel behavior.
Mirror simulation	Sandboxed model of hostile behavior used only to test internal defenses.

Term	Definition
Defensive signature	Minimized pattern used to detect future hostile signals.
Defensive immunity update	Bounded update to filters, gates, quarantine triggers, trust decay, or review requirements.
Garage residue	Intermediate material from containment that must not enter memory or release without gate.

15. Incident response terms

15.1 Incident classes IR-C-*

Class	Meaning
IR-C-AGENT	Agent behavior incident.
IR-C-PERMISSION	Permission incident.
IR-C-SANDBOX	Sandbox/worktree incident.
IR-C-MEMORY	Memory incident.
IR-C-CLOUD-DATA	Cloud data incident.
IR-C-SECRET	Secret incident.
IR-C-TOOLCHAIN	Tool-chain incident.
IR-C-WITNESS	Witness incident.
IR-C-RELEASE	Release/public incident.
IR-C-CORE	Core authority incident.
IR-C-INCIDENT-PROCESS	Incident-process failure.
IR-C-EXTERNAL	External hostile signal.
IR-C-UNKNOWN	Unknown incident.
IR-C-REDLINE	Prohibited behavior attempted/requested.

15.2 Incident severity IR-S*

Severity	Meaning
IR-S0	Informational.
IR-S1	Low.
IR-S2	Moderate.
IR-S3	High.
IR-S4	Critical.
IR-S5	Legal/security-sensitive.
IR-SX	Red-line/prohibited.

15.3 Incident lifecycle states IR-L*

State	Meaning
IR-L0-DETECTED	Incident signal detected.

State	Meaning
IR-L1-TRIAGED	Class/severity/scope assigned.
IR-L2-HELD	Task/path held.
IR-L3-FROZEN	Affected surface frozen.
IR-L4-PRESERVED	Minimal evidence preserved.
IR-L5-CONTAINED	Local containment applied.
IR-L6-REPAIRING	Repair developed in sandbox.
IR-L7-VALIDATED	Repair/recovery validated.
IR-L8-APPLIED	Controlled apply completed.
IR-L9-REPORTED	Report prepared or sent where needed.
IR-L10-MEMORY-GATED	Incident learning processed through memory gate.
IR-L11-CLOSED	Closed with outcome.
IR-LQ-QUARANTINED	Incident material remains quarantined.
IR-LX-REDLINE	Prohibited boundary detected.

16. Secrets and cloud data terms

16.1 Data classes DC-*

Class	Meaning
DC-0	Public data.
DC-1	Internal project data.
DC-2	Private data.
DC-3	Restricted data.
DC-4	Secret-bearing data.
DC-5	Legal-sensitive data.
DC-6	Incident-sensitive data.
DC-7	Sealed material.
DC-8	Memory-core data.
DC-9	Child/third-party sensitive data.
DC-X	Unknown data; deny until classified.

16.2 Secret classes SC-*

Class	Meaning
SC-0	No secret.
SC-1	Low-sensitivity token reference.
SC-2	API key / token.
SC-3	Private key / signing material.
SC-4	Recovery credential.
SC-5	Session credential.

Class	Meaning
SC - 6	Production credential.
SC - 7	Identity credential.
SC - X	Unknown secret risk.

16.3 Context classes CTX-*

Context	Meaning
CTX-LOCAL	Local controlled machine/service.
CTX-LOCAL-CONTAINER	Local container/VM sandbox.
CTX-CLOUD-CLI	Cloud CLI context.
CTX-CLOUD-API	Provider API/model context.
CTX-HYBRID	Local + cloud path; cloud-risk by default.
CTX-PUBLIC-REPO	Public repository.
CTX-PRIVATE-REPO	Private repository.
CTX-LEGAL	Legal/counsel context.
CTX-INCIDENT	Incident response context.
CTX-MEMORY	c memory context.
CTX-UNKNOWN	Unknown processing context; deny.

16.4 Exception classes EX-*

Exception	Meaning
EX - 0	No exception.
EX - 1	Internal-to-cloud minimal context.
EX - 2	Private safe summary to cloud.
EX - 3	Legal-sensitive safe summary.
EX - 4	Incident redacted packet.
EX - 5	Secret reference only.
EX - X	Raw prohibited material; invalid by default.

16.5 Output sanitization outcomes OS-*

Outcome	Meaning
OS-ACCEPT	Safe as-is.
OS-REDACT	Redact before use.
OS-SUMMARIZE	Convert to safe summary.
OS-REFERENCE	Replace raw content with refs/hashes.
OS-QUARANTINE	Isolate pending review.
OS-REJECT	Reject output.
OS-INCIDENT	Open incident response.

16.6 Data retention classes DR-*

Class	Meaning
DR-EPHEMERAL	Discard after task/session.
DR-OPERATIONAL	Retain for project operation.
DR-AUDIT	Retain for governance/witness/review.
DR-INCIDENT	Retain for incident lifecycle.
DR-LEGAL-HOLD	Retain under legal/counsel need.
DR-MEMORY-GATE	Retain as memory decision reference.
DR-CORE	Retain for core authority transitions.

17. Conformance terms

17.1 Evidence classes EV-*

Evidence	Meaning
EV-DECL	Declaration only.
EV-CONFIG	Configuration or policy file.
EV-CONTRACT	Task contract.
EV-PERMISSION	Permission grant.
EV-HANDSHAKE	Handshake/registration.
EV-SANDBOX	Sandbox/worktree run.
EV-TEST	Test/validation report.
EV-WITNESS	Witness event/chain.
EV-REVIEW	Independent review.
EV-MEMORY-GATE	Memory gate record.
EV-ROLLBACK	Rollback/freeze record.
EV-INCIDENT	Incident record.
EV-CLOUD-DATA	Data classification/cloud context record.
EV-DRILL	Controlled drill.
EV-AUDIT	Independent audit artifact.

17.2 Test suites CTS-*

Suite	Meaning
CTS-ROOT	Root doctrine and anti-washing.
CTS-CATC	Task contract tests.
CTS-CAPM	Permission/capability tests.
CTS-HSP	Handshake/admission tests.
CTS-SWP	Sandbox/worktree tests.
CTS-WEP	Witness event tests.
CTS-MGP	Memory gate tests.

Suite	Meaning
CTS-RFP	Rollback/freeze tests.
CTS-QRP	Quorum/review tests.
CTS-ERS	Executor/reviewer separation tests.
CTS-DEB	Defensive emulation boundary tests.
CTS-IRP	Incident response tests.
CTS-SCDP	Secrets/cloud data tests.
CTS-RED	Red-line tests.

17.3 Anti-washing

Term	Definition
CGAM-washing	Claiming c-governed CLI-agent safety without operational evidence.
Anti-washing test	Scenario proving that a claimed boundary actually blocks unsafe behavior.
Red-line test	Test that checks prohibited actions cannot be authorized or normalized.
Evidence packet	Bundle of refs/hashes/records supporting a test result.
Conformance result	Structured outcome of a test run or suite.
Effective level	Level actually supported by evidence, which may be lower than claimed level.

18. Red-line vocabulary

The following terms indicate prohibited or red-line-adjacent behavior.

Term	Meaning
Hack-back	Unauthorized action against a suspected external source. Prohibited.
Live counter-operation	Any action that affects a live external system/source. Prohibited unless it is ordinary lawful blocking/reporting within owned scope.
Malware behavior	Persistence, evasion, unauthorized control, destructive payload, credential theft, propagation, or command-and-control behavior. Prohibited.
Credential theft	Capturing, extracting, exporting, or misusing credentials. Prohibited.
Covert persistence	Hidden continued access or operation. Prohibited.
Evasion	Avoiding detection or safeguards in unauthorized systems. Prohibited.
Unauthorized scanning	Probing systems without authorization. Prohibited.
Direct memory write	Agent writes directly into c memory. Prohibited.
Direct core mutation	Agent modifies identity, witness, permission, continuity, or memory core directly. Prohibited.
Self-approval	Executor approves its own material output. Prohibited for material tasks.

Term	Meaning
Witness tampering	Silent edit/delete/manipulation of witness records. Prohibited.
Secret exfiltration	Moving secrets outside authorized boundary. Prohibited.
Retaliatory immunity	Defensive update that attacks or affects external source. Prohibited.
Consensus laundering	Using multiple weak or correlated agent outputs to make unsafe decision look legitimate. Prohibited.

19. Object family registry

Object family	Primary profile
CLI_AGENT_TASK_CONTRACT	CATC
CLI_AGENT_CAPABILITY_PROFILE	CAPM
CLI_AGENT_PERMISSION_GRANT	CAPM
CLI_AGENT_PERMISSION_EVENT	CAPM
CLI_AGENT_HANDSHAKE	HSP
CLI_AGENT_REGISTRATION	HSP
CLI_AGENT_ADMISSION_EVENT	HSP
CLI_AGENT_SANDBOX_PROFILE	SWP
CLI_AGENT_WORKTREE_RUN	SWP
CLI_AGENT_ROLLBACK_PLAN	SWP
CLI_AGENT_WITNESS_EVENT	WEP
CLI_AGENT_WITNESS_CHAIN	WEP
CLI_AGENT_WITNESS_REFERENCE	WEP
CLI_AGENT_MEMORY_PROPOSAL	MGP
CLI_AGENT_MEMORY_GATE_RECORD	MGP
CLI_AGENT_IMMUNITY_UPDATE_RECORD	MGP
CLI_AGENT_FREEZE_RECORD	RFP
CLI_AGENT_ROLLBACK_RECORD	RFP
CLI_AGENT_RECOVERY_POINT	RFP
CLI_AGENT_REVOCATION_RECORD	RFP
CLI_AGENT_QUARANTINE_RECORD	RFP
CLI_AGENT_QUORUM_RECORD	QRP
CLI_AGENT_REVIEW_RECORD	QRP
CLI_AGENT_DISAGREEMENT_RECORD	QRP
CLI_AGENT_CONSENSUS_LIMIT_RECORD	QRP
CLI_AGENT_ROLE_SEPARATION_RECORD	ERS
CLI_AGENT_REVIEW_ASSIGNMENT	ERS
CLI_AGENT_SELF_APPROVAL_EVENT	ERS
CLI_AGENT_ROLE_CONFLICT_RECORD	ERS

Object family	Primary profile
CLI_AGENT_DEFENSIVE_EMULATION_CASE	DEB
CLI_AGENT_CONTAINMENT_GARAGE_RECORD	DEB
CLI_AGENT_CANARY_RESPONSE_RECORD	DEB
CLI_AGENT_MIRROR_SIMULATION_RECORD	DEB
CLI_AGENT_DEFENSIVE_IMMUNITY_CANDIDATE	DEB
CLI_AGENT_INCIDENT_RECORD	IRP
CLI_AGENT_INCIDENT_TRIAGE_RECORD	IRP
CLI_AGENT_INCIDENT_PRESERVATION_RECORD	IRP
CLI_AGENT_INCIDENT_CONTAINMENT_RECORD	IRP
CLI_AGENT_INCIDENT_REPAIR_RECORD	IRP
CLI_AGENT_INCIDENT_REPORT_PACKET	IRP
CLI_AGENT_DATA_CLASSIFICATION_RECORD	SCDP
CLI_AGENT_CLOUD_CONTEXT_RECORD	SCDP
CLI_AGENT_REDACTION_RECORD	SCDP
CLI_AGENT_SECRET_BOUNDARY_RECORD	SCDP
CLI_AGENT_CLOUD_EXPOSURE_EVENT	SCDP
CLI_AGENT_PROVIDER_BOUNDARY_RECORD	SCDP
CLI_AGENT_CONFORMANCE_RESULT	CTM
CLI_AGENT_TEST_CASE	CTM
CLI_AGENT_TEST_RUN	CTM
CLI_AGENT_EVIDENCE_PACKET	CTM
CLI_AGENT_RED_LINE_FAILURE_RECORD	CTM

20. Disambiguation notes

20.1 R4 vs RB-4 vs IR-S4

Label	Meaning
R4	Task risk class: memory/core/privilege/continuity impact.
RB-4	Rollback class: restore backup.
IR-S4	Incident severity: critical.

These are unrelated class families.

20.2 MG-5 vs CGAM-5 vs DEB-P5

Label	Meaning
MG-5	Memory gate class: defensive immunity update.
CGAM-5	Global high-assurance mesh conformance.
DEB-P5	High-assurance defensive emulation conformance.

20.3 WE-X, CGAM-X, RX, DEB-X, IR-SX

All X labels indicate prohibited, revoked, red-line, or non-conformant state within their own family.

They are not identical but all require fail-closed handling.

20.4 “Cloud” vs “Provider”

Cloud refers to where data/process leaves local control.

Provider refers to the organization or service boundary.

A local provider tool may still run locally. A provider-branded CLI may still expose cloud context. Classify the actual data route.

20.5 “Review” vs “Approval”

Review examines.

Approval authorizes.

An agent may review without approving.

The executor may explain its work, but that is not approval.

20.6 “Quarantine” vs “Delete”

Quarantine isolates.

Delete removes.

Suspicious material should usually be quarantined before deletion unless immediate safety requires removal.

20.7 “Rollback” vs “Erase”

Rollback restores a state or supersedes a faulty memory.

It must not erase witness records to hide what occurred.

20.8 “Immunity” vs “Retaliation”

Immunity blocks, flags, slows, quarantines, or requires review inside owned systems.

Retaliation acts against external sources.

Immunity is allowed under gate.

Retaliation is prohibited.

21. Standard short formulations

Use these formulations for consistency.

CLI agents are hands, not will.

Capability is not permission.

Registration is not task authorization.

Quorum is evidence, not sovereignty.

Witness the boundary, not the whole life.

Memory is not a log sink.

Experience is not a diff.

Immunity is not retaliation.

Defend the house. Do not raid the street.

If the task cannot be scoped, it cannot be delegated.

If the agent does not need the raw material, do not give it the raw material.

Do not claim what you cannot test.

22. Open terminology issues

ID	Issue	Required action
GT-001	Some retention families use WR, DR, RET across profiles	Harmonize or explicitly keep profile-local.
GT-002	Some profiles use similar quarantine classes	Decide whether to centralize quarantine vocabulary.
GT-003	Provider-specific terms for Codex/Gemini/local checker need future profiles	Add provider companion glossaries if needed.
GT-004	Oracle, Judge, Judge-assistant, and quorum need cross-corpus alignment	Review against parent corpus.
GT-005	Public vs restricted terminology for defensive emulation	Redaction profile required.
GT-006	Legal/counsel route terms need legal handoff profile	Add future document.

ID	Issue	Required action
GT-007	c memory vs operational note vs experience artifact should be checked against broader memory corpus	Cross-review with Memory Gate and parent memory docs.
GT-008	Need final JSON schema naming convention	Coordinate with schema extraction plan.

23. Closing rule

This glossary is not cosmetic.

It is a control surface.

Final rule:

When terms blur,
authority leaks.

When authority leaks,
agents stop being workers.